

Importing the necessary libraries

```
In [103]: import pandas as pd
import numpy as np
import math
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import StratifiedKFold
from sklearn.ensemble import VotingClassifier
from sklearn.pipeline import Pipeline
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import style
style.use('ggplot')
```

```
In [104]: !pip install catboost
```

```
Requirement already satisfied: catboost in /usr/local/lib/python3.6/dist-packages (0.24.1)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packages (from catboost) (3.2.2)
Requirement already satisfied: plotly in /usr/local/lib/python3.6/dist-packages (from catboost) (4.4.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from catboost) (1.4.1)
Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.6/dist-packages (from catboost) (1.18.5)
Requirement already satisfied: pandas>=0.24.0 in /usr/local/lib/python3.6/dist-packages (from catboost) (1.1.2)
Requirement already satisfied: graphviz in /usr/local/lib/python3.6/dist-packages (from catboost) (0.10.1)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from catboost) (1.15.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (2.4.7)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (0.10.0)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (1.2.0)
Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.6/dist-packages (from plotly->catboost) (1.3.3)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-packages (from pandas->0.24.0->catboost) (2018.9)
```

```
In [105]: import xgboost as xgb
from xgboost import plot_tree
from catboost import CatBoostClassifier
```

Reading in the train, test and samplesubmission files

```
In [106]: train = pd.read_csv('/Train(1).csv')
test = pd.read_csv('/Test(1).csv')
sub = pd.read_csv('/SampleSubmission(1).csv')
```

```
In [107]: Train = train.copy()
Train = Train.set_index(['Applicant_ID'])
Train.head()
```

```
Out[107]:
```

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form_fiel
Applicant_ID							
Apcnt_1000000	3436.0	0.28505	1.6560	0.0	0.000	0.0	10689720
Apcnt_1000004	3456.0	0.67400	0.2342	0.0	0.000	0.0	898979
Apcnt_1000008	3276.0	0.53845	3.1510	0.0	6.282	NaN	956940
Apcnt_1000012	3372.0	0.17005	0.5050	0.0	0.000	192166.0	3044703
Apcnt_1000016	3370.0	0.77270	1.1010	0.0	0.000	1556.0	214728

Understanding the data set(train and test)

```
In [108]: Train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 56000 entries, Apcnt_1000000 to Apcnt_999996
Data columns (total 51 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   form_field1           53471 non-null    float64
1   form_field2           52156 non-null    float64
2   form_field3           55645 non-null    float64
3   form_field4           55645 non-null    float64
4   form_field5           55645 non-null    float64
5   form_field6           42640 non-null    float64
6   form_field7           50837 non-null    float64
7   form_field8           42640 non-null    float64
8   form_field9           47992 non-null    float64
9   form_field10          55645 non-null    float64
10  form_field11          24579 non-null    float64
11  form_field12          46105 non-null    float64
12  form_field13          50111 non-null    float64
13  form_field14          56000 non-null    int64
14  form_field15          33525 non-null    float64
15  form_field16          42964 non-null    float64
16  form_field17          44849 non-null    float64
17  form_field18          45598 non-null    float64
18  form_field19          55996 non-null    float64
19  form_field20          55645 non-null    float64
20  form_field21          40146 non-null    float64
21  form_field22          35600 non-null    float64
22  form_field23          27877 non-null    float64
23  form_field24          42703 non-null    float64
24  form_field25          50550 non-null    float64
25  form_field26          48562 non-null    float64
26  form_field27          46701 non-null    float64
27  form_field28          55645 non-null    float64
28  form_field29          55645 non-null    float64
29  form_field30          30491 non-null    float64
30  form_field31          16592 non-null    float64
31  form_field32          50550 non-null    float64
32  form_field33          54744 non-null    float64
33  form_field34          55645 non-null    float64
34  form_field35          32852 non-null    float64
35  form_field36          54005 non-null    float64
36  form_field37          50550 non-null    float64
37  form_field38          55645 non-null    float64
38  form_field39          51789 non-null    float64
39  form_field40          12271 non-null    float64
40  form_field41          17771 non-null    float64
41  form_field42          54677 non-null    float64
42  form_field43          55432 non-null    float64
43  form_field44          50617 non-null    float64
44  form_field45          24683 non-null    float64
45  form_field46          40096 non-null    float64
46  form_field47          56000 non-null    object
47  form_field48          35111 non-null    float64
48  form_field49          55645 non-null    float64
49  form_field50          44944 non-null    float64
```

```

50 default_status 56000 non-null object
dtypes: float64(48), int64(1), object(2)
memory usage: 22.2+ MB

```

In [109]: Train.describe()

Out[109]:

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	forr
count	53471.000000	52156.000000	55645.000000	55645.000000	55645.000000	4.264000e+04	5.0837
mean	3491.795665	0.550737	1.052225	0.851979	1.956317	6.244479e+05	6.8652
std	188.462426	0.820979	2.147768	3.157692	10.512396	1.433422e+06	1.9127
min	2990.000000	0.000000	0.000000	0.000000	0.000000	0.000000e+00	0.0000
25%	3358.000000	0.070788	0.000000	0.000000	0.000000	1.400400e+04	6.8697
50%	3484.000000	0.267575	0.062000	0.000000	0.000000	1.155330e+05	2.7043
75%	3620.000000	0.719512	1.282000	0.000000	0.000000	5.259280e+05	6.9938
max	3900.000000	18.015050	57.371600	91.672200	407.748600	5.313546e+07	2.1587

In [110]: Train.describe(include = 'all')

Out[110]:

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	for
count	53471.000000	52156.000000	55645.000000	55645.000000	55645.000000	4.264000e+04	5.0837
unique	NaN	NaN	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	NaN	NaN	
mean	3491.795665	0.550737	1.052225	0.851979	1.956317	6.244479e+05	6.8652
std	188.462426	0.820979	2.147768	3.157692	10.512396	1.433422e+06	1.9127
min	2990.000000	0.000000	0.000000	0.000000	0.000000	0.000000e+00	0.0000
25%	3358.000000	0.070788	0.000000	0.000000	0.000000	1.400400e+04	6.8697
50%	3484.000000	0.267575	0.062000	0.000000	0.000000	1.155330e+05	2.7043
75%	3620.000000	0.719512	1.282000	0.000000	0.000000	5.259280e+05	6.9938
max	3900.000000	18.015050	57.371600	91.672200	407.748600	5.313546e+07	2.1587

Finding numeric and categorical columns

```
In [111]: numeric = Train.select_dtypes(exclude = 'object')
numeric.columns
```

```
Out[111]: Index(['form_field1', 'form_field2', 'form_field3', 'form_field4',
                'form_field5', 'form_field6', 'form_field7', 'form_field8',
                'form_field9', 'form_field10', 'form_field11', 'form_field12',
                'form_field13', 'form_field14', 'form_field15', 'form_field16',
                'form_field17', 'form_field18', 'form_field19', 'form_field20',
                'form_field21', 'form_field22', 'form_field23', 'form_field24',
                'form_field25', 'form_field26', 'form_field27', 'form_field28',
                'form_field29', 'form_field30', 'form_field31', 'form_field32',
                'form_field33', 'form_field34', 'form_field35', 'form_field36',
                'form_field37', 'form_field38', 'form_field39', 'form_field40',
                'form_field41', 'form_field42', 'form_field43', 'form_field44',
                'form_field45', 'form_field46', 'form_field48', 'form_field49',
                'form_field50'],
                dtype='object')
```

```
In [112]: objects = Train.select_dtypes(include = 'object')
objects.columns
```

```
Out[112]: Index(['form_field47', 'default_status'], dtype='object')
```

```
In [113]: Train.isnull().sum()
```

```
Out[113]: form_field1      2529
form_field2      3844
form_field3       355
form_field4       355
form_field5       355
form_field6     13360
form_field7       5163
form_field8     13360
form_field9       8008
form_field10      355
form_field11     31421
form_field12       9895
form_field13       5889
form_field14         0
form_field15     22475
form_field16     13036
form_field17     11151
form_field18     10402
form_field19         4
form_field20         4
```

```
In [114]: Train.shape
```

```
Out[114]: (56000, 51)
```

```
In [115]: Train['form_field47'].value_counts()
```

```
Out[115]: charge      36373
lending      19627
Name: form_field47, dtype: int64
```

```
In [116]: Test = test.set_index(['Applicant_ID'])
Test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 24000 entries, Apcnt_1000032 to Apcnt_999992
Data columns (total 50 columns):
#   Column          Non-Null Count  Dtype
---  -
0   form_field1     22890 non-null    float64
1   form_field2     22291 non-null    float64
2   form_field3     23854 non-null    float64
3   form_field4     23854 non-null    float64
4   form_field5     23854 non-null    float64
5   form_field6     18396 non-null    float64
6   form_field7     21769 non-null    float64
7   form_field8     18396 non-null    float64
8   form_field9     20600 non-null    float64
9   form_field10    23853 non-null    float64
10  form_field11    10602 non-null    float64
11  form_field12    19817 non-null    float64
12  form_field13    21537 non-null    float64
13  form_field14    24000 non-null    int64
14  form_field15    14408 non-null    float64
15  form_field16    18526 non-null    float64
16  form_field17    19305 non-null    float64
17  form_field18    19631 non-null    float64
18  form_field19    24000 non-null    float64
19  form_field20    23853 non-null    float64
20  form_field21    17293 non-null    float64
21  form_field22    15276 non-null    float64
22  form_field23    11875 non-null    float64
23  form_field24    18395 non-null    float64
24  form_field25    21744 non-null    float64
25  form_field26    20828 non-null    float64
26  form_field27    20090 non-null    float64
27  form_field28    23853 non-null    float64
28  form_field29    23853 non-null    float64
29  form_field30    13092 non-null    float64
30  form_field31    7190 non-null     float64
31  form_field32    21744 non-null    float64
32  form_field33    23505 non-null    float64
33  form_field34    23853 non-null    float64
34  form_field35    14134 non-null    float64
35  form_field36    23097 non-null    float64
36  form_field37    21744 non-null    float64
37  form_field38    23853 non-null    float64
38  form_field39    22171 non-null    float64
39  form_field40    5172 non-null     float64
40  form_field41    7651 non-null     float64
41  form_field42    23422 non-null    float64
42  form_field43    23750 non-null    float64
43  form_field44    21638 non-null    float64
44  form_field45    10462 non-null    float64
45  form_field46    17115 non-null    float64
46  form_field47    24000 non-null    object
47  form_field48    15078 non-null    float64
```

```

48 form_field49 23854 non-null float64
49 form_field50 19203 non-null float64
dtypes: float64(48), int64(1), object(1)
memory usage: 9.3+ MB

```

In [117]: Test.isnull().sum()

```

Out[117]: form_field1      1110
form_field2      1709
form_field3       146
form_field4       146
form_field5       146
form_field6      5604
form_field7      2231
form_field8      5604
form_field9      3400
form_field10      147
form_field11     13398
form_field12      4183
form_field13      2463
form_field14        0
form_field15      9592
form_field16      5474
form_field17      4695
form_field18      4369
form_field19        0
form_field20      147

```

In [118]: Test['form_field47'].value_counts()

```

Out[118]: charge      15467
lending       8533
Name: form_field47, dtype: int64

```

Preprocessing and cleaning of the data set(train and test)

In [119]: Train.replace('no', 0, inplace = True)
Train.replace('yes', 1, inplace = True)
Train.head(3)

```

Out[119]:
```

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form_fiel
Applicant_ID							
Apcnt_1000000	3436.0	0.28505	1.6560	0.0	0.000	0.0	10689720
Apcnt_1000004	3456.0	0.67400	0.2342	0.0	0.000	0.0	898979
Apcnt_1000008	3276.0	0.53845	3.1510	0.0	6.282	NaN	956940

```
In [120]: Train = pd.get_dummies(Train)
Train.head()
```

```
Out[120]:
```

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form_fiel
Applicant_ID							
Apcnt_1000000	3436.0	0.28505	1.6560	0.0	0.000	0.0	10689720
Apcnt_1000004	3456.0	0.67400	0.2342	0.0	0.000	0.0	898979
Apcnt_1000008	3276.0	0.53845	3.1510	0.0	6.282	NaN	956940
Apcnt_1000012	3372.0	0.17005	0.5050	0.0	0.000	192166.0	3044703
Apcnt_1000016	3370.0	0.77270	1.1010	0.0	0.000	1556.0	214728

```
In [121]: Test = pd.get_dummies(Test)
Test.head(3)
```

```
Out[121]:
```

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form_fiel
Applicant_ID							
Apcnt_1000032	3236.0	0.34875	10.2006	0.0000	0.0	418564.0	418564
Apcnt_1000048	3284.0	1.27360	2.9606	9.0198	0.0	0.0	9858816
Apcnt_1000052	NaN	0.27505	0.0600	0.0000	0.0	NaN	NaN

```
In [122]: Train.fillna(-999, inplace = True)
Train.head()
```

```
Out[122]:
```

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form_fiel
Applicant_ID							
Apcnt_1000000	3436.0	0.28505	1.6560	0.0	0.000	0.0	10689720
Apcnt_1000004	3456.0	0.67400	0.2342	0.0	0.000	0.0	898979
Apcnt_1000008	3276.0	0.53845	3.1510	0.0	6.282	-999.0	956940
Apcnt_1000012	3372.0	0.17005	0.5050	0.0	0.000	192166.0	3044703
Apcnt_1000016	3370.0	0.77270	1.1010	0.0	0.000	1556.0	214728


```
In [123]: Test.fillna(-999, inplace = True)
Test.head()
```

Out[123]:

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form_fiel
Applicant_ID							
Apcnt_1000032	3236.0	0.34875	10.2006	0.0000	0.0	418564.0	418564
Apcnt_1000048	3284.0	1.27360	2.9606	9.0198	0.0	0.0	9858816
Apcnt_1000052	-999.0	0.27505	0.0600	0.0000	0.0	-999.0	-999
Apcnt_1000076	3232.0	0.28505	2.8032	0.0000	0.0	0.0	473802
Apcnt_1000080	3466.0	2.09545	0.8318	2.5182	0.0	19839.0	1150662



Visualizing and Exploring the train set


```
In [124]: labels = 'default', 'undefault'
default_size = 0
undefault_size = 0
for i in Train['default_status']:
    if i == 1:
        default_size += 1
    else:
        undefault_size +=1
print(default_size)
print(undefault_size)
sizes = [(default_size / len(Train)) * 100, (undefault_size / len(Train)) * 100]
explode = (0.1, 0) # only "explode" the 1st slice (i.e. 'default')

fig1, ax1 = plt.subplots()
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%',
        shadow=True, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.


plt.show()
```

13715

42285



```
In [125]: Train.hist(bins=10, figsize=(20, 15))
plt.show()
```




```
In [126]: sns.distplot(Train['form_field29'], bins=10)
```

/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[126]: <matplotlib.axes._subplots.AxesSubplot at 0x7f496c608240>
```



Determining correlation between target and features

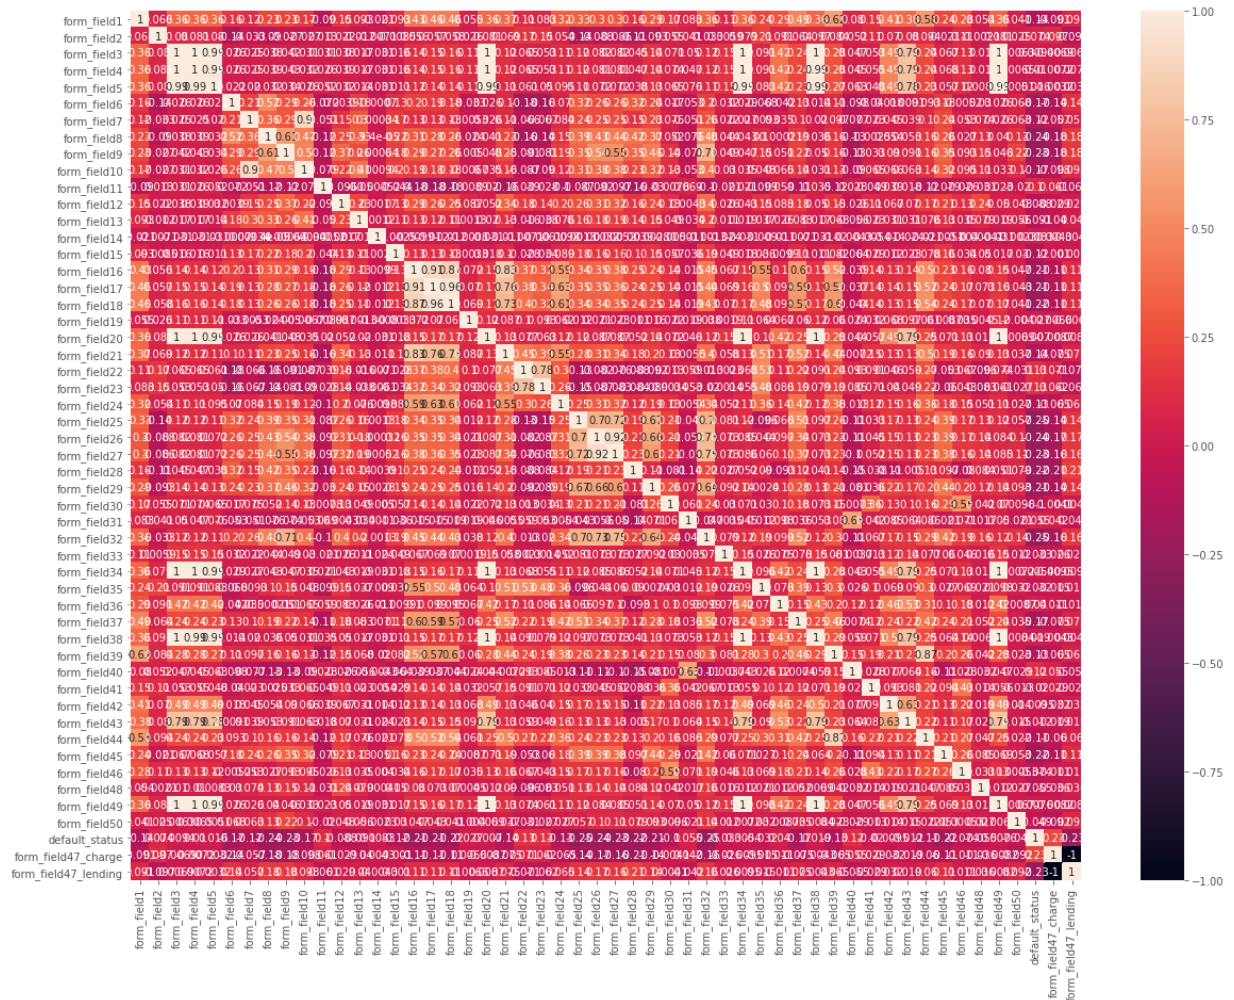
```
In [127]: corrMatrix = Train.corr()
corrMatrix
```

```
Out[127]:
```

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form
form_field1	1.000000	0.066395	0.358717	0.359764	0.355599	0.155828	(
form_field2	0.066395	1.000000	0.080174	0.080579	0.080074	-0.138547	-(
form_field3	0.358717	0.080174	1.000000	0.999254	0.993343	0.025861	(
form_field4	0.359764	0.080579	0.999254	1.000000	0.991149	0.026420	(
form_field5	0.355599	0.080074	0.993343	0.991149	1.000000	0.021725	(
form_field6	0.155828	-0.138547	0.025861	0.026420	0.021725	1.000000	(
form_field7	0.123642	-0.032814	0.024715	0.025235	0.020352	0.207660	1
form_field8	0.230769	-0.090430	0.038259	0.039185	0.032050	0.522087	(
form_field9	0.229502	-0.026720	0.041953	0.042613	0.034483	0.289385	(
form_field10	0.172198	-0.027382	0.031477	0.032431	0.025820	0.258334	(
form_field11	-0.089929	0.012688	0.030722	0.025518	0.052391	-0.071671	-(
form_field12	0.150727	0.022158	0.038149	0.038887	0.031910	0.003930	(
form_field13	0.093389	-0.012013	0.016790	0.017195	0.013911	0.183866	(
form_field14	-0.021031	-0.007145	-0.030710	-0.030709	-0.030823	-0.000700	(
form_field15	0.092517	-0.008544	0.015772	0.016036	0.011432	0.132512	(
form_field16	0.431777	0.055683	0.136770	0.137795	0.121704	0.204917	(
form_field17	0.458290	0.057080	0.152245	0.152967	0.135747	0.185440	(
form_field18	0.460704	0.057772	0.159336	0.159954	0.142585	0.177611	(
form_field19	0.054589	0.026108	0.111975	0.111989	0.110062	-0.033103	-(
form_field20	0.361945	0.081049	0.999346	0.998931	0.991007	0.025788	(
form_field21	0.371567	0.069337	0.119697	0.120604	0.106645	0.104013	(
form_field22	0.108731	0.166758	0.065471	0.065385	0.061349	-0.178472	-(
form_field23	0.083393	0.145292	0.053084	0.052822	0.049722	-0.162650	-(
form_field24	0.322807	0.054365	0.107872	0.106595	0.094787	0.069625	(
form_field25	0.332689	-0.140681	0.118576	0.117680	0.108797	0.318613	(
form_field26	0.302250	-0.088207	0.082369	0.081276	0.072006	0.255407	(
form_field27	0.303718	-0.086130	0.082392	0.081334	0.071907	0.255463	(
form_field28	0.155923	-0.113445	0.045485	0.047075	0.037931	0.320936	(
form_field29	0.291305	-0.092637	0.140925	0.140412	0.134865	0.237581	(
form_field30	0.171175	0.055030	0.070512	0.073910	0.064811	0.017363	(
form_field31	0.082676	0.041341	0.049570	0.046715	0.075906	-0.052720	-(
form_field32	0.355534	-0.032589	0.118484	0.117139	0.106186	0.199499	(
form_field33	0.106149	0.005889	0.150281	0.150386	0.147800	0.031733	(
form_field34	0.360407	0.079351	0.999621	0.999212	0.991395	0.028987	(

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form
form_field35	0.241815	0.206719	0.091056	0.091366	0.082693	-0.067531	(
form_field36	0.285131	0.091039	0.417024	0.416246	0.416548	-0.042113	(
form_field37	0.491966	0.063798	0.241281	0.240333	0.229409	0.125746	(
form_field38	0.364180	0.097107	0.995056	0.994600	0.986949	0.013946	(
form_field39	0.622855	0.083604	0.276831	0.277598	0.268839	0.106078	(
form_field40	0.079766	0.052440	0.046772	0.045198	0.062885	-0.098012	-(
form_field41	0.150026	0.113881	0.052861	0.054992	0.048353	-0.039541	(
form_field42	0.413058	0.070366	0.491087	0.490614	0.488447	0.018392	(
form_field43	0.381166	0.080117	0.787299	0.786950	0.780580	0.009114	(
form_field44	0.579925	0.094005	0.241327	0.242454	0.232281	0.092688	(
form_field45	0.240394	-0.021014	0.067352	0.067842	0.056543	0.179975	(
form_field46	0.281385	0.105363	0.127518	0.133351	0.121485	-0.005218	(
form_field48	0.054497	-0.002065	0.010087	0.010339	0.008272	0.030454	(
form_field49	0.361472	0.080659	0.999499	0.998938	0.992554	0.025554	(
form_field50	0.040527	0.024883	0.006288	0.006456	0.005129	0.067747	(
default_status	-0.140965	0.073846	0.009359	0.009983	0.015713	-0.173197	-(
form_field47_charge	-0.090533	0.097002	-0.006916	-0.007203	-0.003183	-0.139724	-(
form_field47_lending	0.090533	-0.097002	0.006916	0.007203	0.003183	0.139724	(

```
In [128]: plt.subplots(figsize = (20, 15))
sns.heatmap(corrMatrix, annot = True)
plt.show()
```



In [129]: Train.columns

```
Out[129]: Index(['form_field1', 'form_field2', 'form_field3', 'form_field4',
                'form_field5', 'form_field6', 'form_field7', 'form_field8',
                'form_field9', 'form_field10', 'form_field11', 'form_field12',
                'form_field13', 'form_field14', 'form_field15', 'form_field16',
                'form_field17', 'form_field18', 'form_field19', 'form_field20',
                'form_field21', 'form_field22', 'form_field23', 'form_field24',
                'form_field25', 'form_field26', 'form_field27', 'form_field28',
                'form_field29', 'form_field30', 'form_field31', 'form_field32',
                'form_field33', 'form_field34', 'form_field35', 'form_field36',
                'form_field37', 'form_field38', 'form_field39', 'form_field40',
                'form_field41', 'form_field42', 'form_field43', 'form_field44',
                'form_field45', 'form_field46', 'form_field48', 'form_field49',
                'form_field50', 'default_status', 'form_field47_charge',
                'form_field47_lending'],
              dtype='object')
```

Determining and Building the best model

```
In [130]: X = np.array(Train.drop(['default_status'], 1))
          y = np.array(Train['default_status'])
```

```
In [131]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random
```

```
In [132]: print(Train.shape)
          print(Test.shape)
```

```
(56000, 52)
(24000, 51)
```

xgboost classifier

```
In [133]: xgb_clf = xgb.XGBClassifier(n_jobs = -1)
          xgb_clf.fit(X_train, y_train)
          preds1 = xgb_clf.predict_proba(X_test)[:,:1]
          accuracy1 = xgb_clf.score(X_test, y_test)
          print(accuracy1)
```

```
0.808125
```

```
In [134]: auc_df = roc_auc_score(y_test, preds1)
          auc_df
```

```
Out[134]: 0.8310124691046903
```

catboost classifier


```
In [135]: cat_clf = CatBoostClassifier(verbose=False)
cat_clf.fit(X_train, y_train)
preds3 = cat_clf.predict_proba(X_test)
accuracy3 = cat_clf.score(X_test, y_test)
print(accuracy3)
```

0.8091071428571428

```
In [136]: auc_cat = roc_auc_score(y_test, preds3[:,1])
auc_cat
```

Out[136]: 0.8370318115528094

lightgbm classifier with manual parameter tuning

```
In [137]: import lightgbm as lgb
lg_clf = lgb.LGBMClassifier(learning_rate=0.01, num_iterations=2500, scale_pos_we
lg_clf.fit(X_train, y_train)
preds4 = lg_clf.predict_proba(X_test)
accuracy4 = lg_clf.score(X_test, y_test)
print(accuracy4)
```

/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))

0.746875

```
In [138]: lg_cat = roc_auc_score(y_test, preds4[:,1])
lg_cat
```

Out[138]: 0.8362065544341879

randomforest classifier

```
In [139]: from sklearn.ensemble import RandomForestClassifier
rf_clf = RandomForestClassifier()
rf_clf.fit(X_train, y_train)
preds5 = rf_clf.predict_proba(X_test)
accuracy5 = rf_clf.score(X_test, y_test)
print(accuracy5)
```

0.8046428571428571

```
In [140]: rf_auc = roc_auc_score(y_test, preds5[:,1])
rf_auc
```

Out[140]: 0.8304646206478884

xgbclassifier tuned manually

```
In [141]: xgb_clf = xgb.XGBClassifier(n_jobs = -1, scale_pos_weight=3, eta=0.01, n_estimators=1000)
xgb_clf.fit(X_train, y_train)
predstun = xgb_clf.predict_proba(X_test)
accuracytun = xgb_clf.score(X_test, y_test)
print(accuracytun)
```

0.7473214285714286

```
In [142]: auc_df = roc_auc_score(y_test, predstun[:,1])
auc_df
```

Out[142]: 0.8353853787702872

Looking for the best model by predicting on the test set using stratified kfold.

```
In [143]: X = Train.drop(['default_status'], 1)
y = Train['default_status']
```

```
In [144]: X.shape
```

Out[144]: (56000, 51)

```
In [145]: Test.shape
```

Out[145]: (24000, 51)

stratified kfold to predict on test data with catboost

```
In [146]: #Test = Test.set_index(['Applicant_ID'])
```

```
In [147]: skfold = StratifiedKFold(n_splits=25, random_state=0, shuffle=True)

scores = []
preds = []
i = 1
cat = CatBoostClassifier(verbose=False)

for train_split, test_split in skfold.split(X, y):
    Xtrain, Xtest, ytrain, ytest = X.iloc[train_split], X.iloc[test_split], y[train_split], y[test_split]
    cat.fit(Xtrain, ytrain)
    score = roc_auc_score(ytest, cat.predict_proba(Xtest)[:,-1])
    pred = cat.predict_proba(Xtest)[:,-1]
    scores.append(score)
    preds.append(pred)
    print('AUC Score for {} split:'.format(i), score)
    i+=1
print('Final score:', np.mean(scores))
Final_pred1 = np.mean(preds, axis = 0)
```

```
AUC Score for 1 split: 0.8325093613570085
AUC Score for 2 split: 0.8460747010405343
AUC Score for 3 split: 0.8353037479939949
AUC Score for 4 split: 0.8322645424582836
AUC Score for 5 split: 0.838814256872185
AUC Score for 6 split: 0.854275594899139
AUC Score for 7 split: 0.845760858311332
AUC Score for 8 split: 0.8234780245379718
AUC Score for 9 split: 0.8328738934617177
AUC Score for 10 split: 0.8365418629531844
AUC Score for 11 split: 0.847058088519635
AUC Score for 12 split: 0.8408471291817066
AUC Score for 13 split: 0.8301152894516022
AUC Score for 14 split: 0.8536191279451161
AUC Score for 15 split: 0.8484519458528436
AUC Score for 16 split: 0.8450243925033312
AUC Score for 17 split: 0.8364641264855514
AUC Score for 18 split: 0.8277250503307448
AUC Score for 19 split: 0.8424251824994425
AUC Score for 20 split: 0.8522328107984087
AUC Score for 21 split: 0.8429325293340185
AUC Score for 22 split: 0.8224178362034514
AUC Score for 23 split: 0.8517405443368352
AUC Score for 24 split: 0.8389847031159282
AUC Score for 25 split: 0.8515822004203116
Final score: 0.840380712034571
```

stratified kfold to predict test data with voting classifier(xgboost and catboost)

```

In [148]: skfold = StratifiedKFold(n_splits=25, random_state=0, shuffle=True)

scores = []
preds = []
i = 1
vc = VotingClassifier(estimators= [('xg', xgb.XGBClassifier(n_jobs = -1, scale_po

for train_split, test_split in skfold.split(X, y):
    Xtrain, Xtest, ytrain, ytest = X.iloc[train_split], X.iloc[test_split], y[train_split], y[test_split]
    vc.fit(Xtrain, ytrain)
    score = roc_auc_score(ytest, vc.predict_proba(Xtest)[:,-1])
    pred = vc.predict_proba(Xtest)[:,-1]
    scores.append(score)
    preds.append(pred)
    print('AUC Score for {} split:'.format(i), score)
    i+=1
print('Final score:', np.mean(scores))
Final_pred2 = np.mean(preds, axis = 0)

```

```

AUC Score for 1 split: 0.8305842435850977
AUC Score for 2 split: 0.8461189194319338
AUC Score for 3 split: 0.8376020258494244
AUC Score for 4 split: 0.831605580576694
AUC Score for 5 split: 0.8394074304153509
AUC Score for 6 split: 0.8538118410381184
AUC Score for 7 split: 0.846544925885662
AUC Score for 8 split: 0.821440743041535
AUC Score for 9 split: 0.831855791962175
AUC Score for 10 split: 0.8372288657659056
AUC Score for 11 split: 0.8459227518664654
AUC Score for 12 split: 0.8419016781223643
AUC Score for 13 split: 0.8306409481676809
AUC Score for 14 split: 0.8521994185439039
AUC Score for 15 split: 0.8476699207957268
AUC Score for 16 split: 0.8449317559263172
AUC Score for 17 split: 0.8358048987514529
AUC Score for 18 split: 0.832837296778509
AUC Score for 19 split: 0.8422054399214097
AUC Score for 20 split: 0.8538841116421557
AUC Score for 21 split: 0.843866435290658
AUC Score for 22 split: 0.8211435446847609
AUC Score for 23 split: 0.8526098201234653
AUC Score for 24 split: 0.8402762293466213
AUC Score for 25 split: 0.8509735996527206
Final score: 0.8405227286866445

```

Stratified kfold to predict test data with voting classifier(lightgbm, catboost and randomforest).

```
In [149]: skfold = StratifiedKFold(n_splits=25, random_state=0, shuffle=True)

scores = []
preds = []
i = 1
vc = VotingClassifier(estimators= [('lg_clf', lgb.LGBMClassifier(learning_rate=0.

for train_split, test_split in skfold.split(X, y):
    Xtrain, Xtest, ytrain, ytest = X.iloc[train_split], X.iloc[test_split], y[train_split], y[test_split]
    vc.fit(Xtrain, ytrain)
    score = roc_auc_score(ytest, vc.predict_proba(Xtest)[:,-1])
    pred = vc.predict_proba(Xtest)[:,-1]
    scores.append(score)
    preds.append(pred)
    print('AUC Score for {} split:'.format(i), score)
    i+=1
print('Final score:', np.mean(scores))
Final_pred3 = np.mean(preds, axis = 0)
```

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 1 split: 0.8340958309606391

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 2 split: 0.8476902900726476

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 3 split: 0.8379741074355922

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 4 split: 0.8334951079360148

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 5 split: 0.8415622681230694

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
```

```
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 6 split: 0.8556722489689566

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 7 split: 0.8477603924004763

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 8 split: 0.8228751445186451

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 9 split: 0.8326851564252558

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 10 split: 0.8385672809787579

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 11 split: 0.848005997679777

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 12 split: 0.8428420470960049

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 13 split: 0.8314068156822954

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 14 split: 0.8563055886785178

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 15 split: 0.8466519956180745

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 16 split: 0.8441820459542052

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 17 split: 0.8371352030841517

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 18 split: 0.8350142563383346

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 19 split: 0.8413943312877885

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 20 split: 0.8532054948570542

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 21 split: 0.8454897297274008

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 22 split: 0.8230113565980401

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 23 split: 0.852983597940021

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument  
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 24 split: 0.8395965353920196

```
/usr/local/lib/python3.6/dist-packages/lightgbm/engine.py:118: UserWarning: Found `num_iterations` in params. Will use it instead of argument  
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
```

AUC Score for 25 split: 0.8548007828867927

Final score: 0.8417761442656212

After several trials, the above model combination gave the best auc score

Catboost model prediction

```
In [150]: default_status1 = Final_pred1
```

Votingclassifier(xgboost and catboost) prediction

```
In [151]: default_status2 = Final_pred2
```

Votingclassifier(lightgbm, catboost and randomforest) prediction

```
In [152]: default_status3 = Final_pred3
```

Weighted average of catboost model and votingclassifier(with xgboost and catboost)

```
In [153]: Final_pred4 = Final_pred1*0.4 + Final_pred2*0.6
```

Weighted average of votingclassifier(with lightgbm, catboost and randomforest) and voting classifier(with xgboost and catboost)

```
In [154]: Final_pred5 = Final_pred3*0.7 + Final_pred2*0.3
```

Weighted average of votingclassifier(with lightgbm, catboost and randomforest) and weighted average of catboost model and votingclassifier(with xgboost and catboost)

```
In [195]: Final_pred6 = Final_pred3*0.7 + Final_pred4*0.3
```



```
In [196]: default_status4 = Final_pred4
          default_status5 = Final_pred5
          default_status6 = Final_pred6
```

```
In [197]: Test = Test.reset_index()
```

Predicted values for votingclassifier with lightgbm, catboost and randomforest.

```
In [198]: predicted_values3 = pd.DataFrame({'Applicant_ID': Test['Applicant_ID'], 'default_
          predicted_values3
```

```
Out[198]:
```

	Applicant_ID	default_status
0	Apcnt_1000032	0.422178
1	Apcnt_1000048	0.404793
2	Apcnt_1000052	0.495254
3	Apcnt_1000076	0.776757
4	Apcnt_1000080	0.229802
...
23995	Apcnt_999940	0.736293
23996	Apcnt_999956	0.356947
23997	Apcnt_999976	0.373381
23998	Apcnt_999984	0.615068
23999	Apcnt_999992	0.276300

24000 rows × 2 columns

Predicted values for weighted average of votingclassifier(with lightgbm, catboost and randomforest) and votingclassifier(with catboost and xgboost).

```
In [199]: predicted_values5 = pd.DataFrame({'Applicant_ID': Test['Applicant_ID'], 'default_
predicted_values5
```

```
Out[199]:
```

	Applicant_ID	default_status
0	Apcnt_1000032	0.429166
1	Apcnt_1000048	0.417463
2	Apcnt_1000052	0.511428
3	Apcnt_1000076	0.792466
4	Apcnt_1000080	0.237686
...
23995	Apcnt_999940	0.754724
23996	Apcnt_999956	0.367176
23997	Apcnt_999976	0.376299
23998	Apcnt_999984	0.617652
23999	Apcnt_999992	0.277152

24000 rows × 2 columns

Predicted values for weighted average of votingclassifier(with lightgbm, catboost and randomforest) and the weighted average of catboost model and votingclassifier(with catboost and xgboost). This gave me my best auc score on the leaderboard.

```
In [200]: predicted_values6 = pd.DataFrame({'Applicant_ID': Test['Applicant_ID'], 'default_
predicted_values6
```

```
Out[200]:
```

	Applicant_ID	default_status
0	Apcnt_1000032	0.411899
1	Apcnt_1000048	0.399594
2	Apcnt_1000052	0.497011
3	Apcnt_1000076	0.783711
4	Apcnt_1000080	0.227037
...
23995	Apcnt_999940	0.748793
23996	Apcnt_999956	0.354005
23997	Apcnt_999976	0.359800
23998	Apcnt_999984	0.606497
23999	Apcnt_999992	0.267104

24000 rows × 2 columns

```
In [201]: predicted_values3.to_csv('FinSub1.csv', index = False)
```

```
In [202]: predicted_values5.to_csv('FinSub2.csv', index = False)
```

The best and final Submission file

```
In [203]: predicted_values6.to_csv('FinSub3.csv', index = False)
```